

Data in Azure Machine Learning

Article • 08/19/2022 • 8 minutes to read

Select the version of Azure Machine Learning developer platform you are using:

Azure Machine Learning lets you bring data from a local machine or an existing cloud-based storage. In this article you will learn the main data concepts in Azure Machine Learning, including:

- ✓ **URIs** - A **Uniform Resource Identifier** that is a reference to a storage location on your local computer or in the cloud that makes it very easy to access data in your jobs.
- ✓ **Data asset** - Create data assets in your workspace to share with team members, version, and track data lineage.
- ✓ **Datastore** - Azure Machine Learning Datastores securely keep the connection information to your data storage on Azure, so you don't have to code it in your scripts.
- ✓ **MLTable** - a method to abstract the schema definition for tabular data so that it is easier for consumers of the data to materialize the table into a Pandas/Dask/Spark dataframe.

URIs

A URI (uniform resource identifier) represents a storage location on your local computer, an attached Datastore, blob/ADLS storage, or a publicly available http(s) location. In addition to local paths (for example: `./path_to_my_data/`), several different protocols are supported for cloud storage locations:

- `http(s)` - Private/Public Azure Blob Storage Locations, or publicly available http(s) location
- `abfs(s)` - Azure Data Lake Storage Gen2 storage location
- `azureml` - An Azure Machine Learning **Datastore** location

Azure Machine Learning distinguishes two types of URIs:

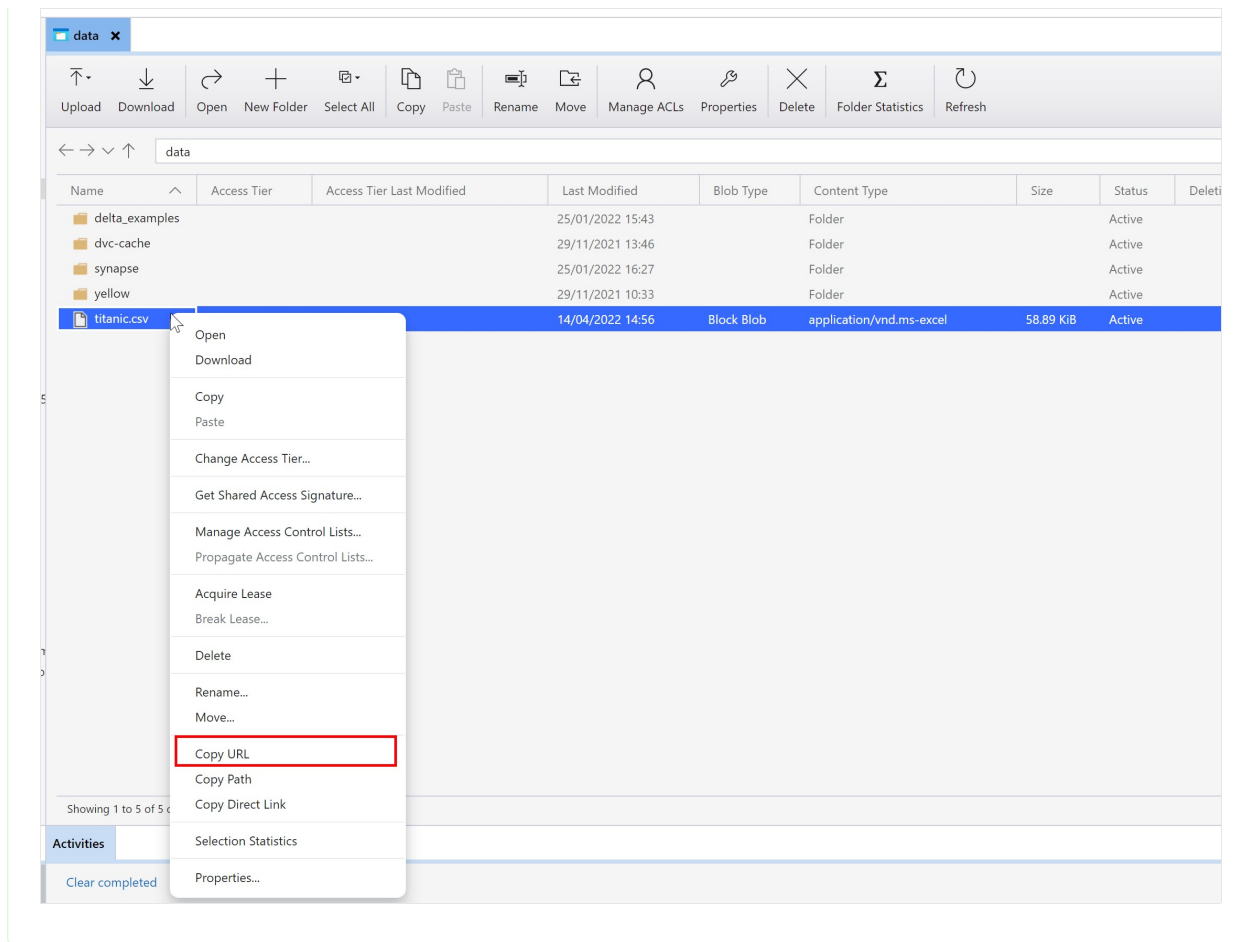
Data type	Description	Examples
-----------	-------------	----------

Data type	Description	Examples
uri_file	Refers to a specific file location	<code>https://<account_name>.blob.core.windows.net/<container_name>/<folder>/<file></code> <code>azureml://datastores/<datastore_name>/paths/<folder>/<file></code> <code>abfss://<file_system>@<account_name>.dfs.core.windows.net/<folder>/<file></code>
uri_folder	Refers to a specific folder location	<code>https://<account_name>.blob.core.windows.net/<container_name>/<folder></code> <code>azureml://datastores/<datastore_name>/paths/<folder></code> <code>abfss://<file_system>@<account_name>.dfs.core.windows.net/<folder>/</code>

URIs are mapped to the filesystem on the compute target, hence using URIs is like using files or folders in the command that consumes/produces them. URIs leverage **identity-based authentication** to connect to storage services with either your Azure Active Directory ID (default) or Managed Identity.

Tip

For data located in an Azure storage account we recommend using the **Azure Storage Explorer** . You can browse data and obtain the URI for any file/folder by right-selecting **Copy URL**:



Examples

uri_file

Below is an example of a job specification that shows how to access a file from a public blob store. In this example, the job executes the Linux `ls` command.

yml

```
# hello-data-uri-file.yml
$schema: https://azuremlschemas.azureedge.net/latest/commandJob.schema.json
command: |
  ls ${inputs.my_csv_file}

inputs:
  my_csv_file:
    type: uri_file
    path: https://azuremlexamples.blob.core.windows.net/datasets/titanic.csv
environment: azureml:AzureML-sklearn-1.0-ubuntu20.04-py38-cpu@latest
```

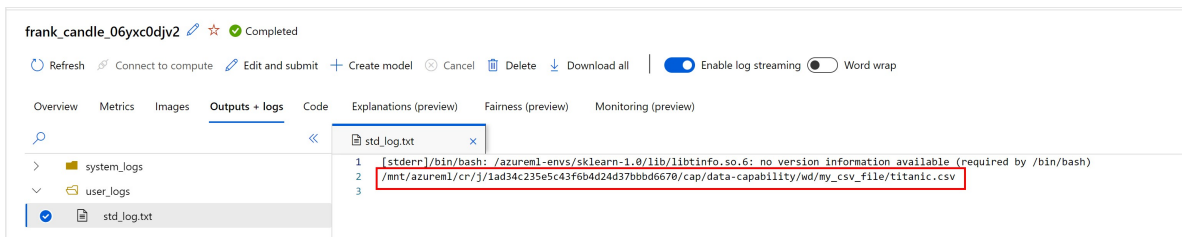
```
compute: azureml:cpu-cluster
```

Create the job using the CLI:

Azure CLI

```
az ml job create --file hello-data-uri-file.yml
```

When the job has completed the user logs will show the standard output of the Linux command `ls ${inputs.my_csv_file}`:



Notice that the file has been mapped to the filesystem on the compute target and `${inputs.my_csv_file}` resolves to that location.

Data asset

Azure Machine Learning allows you to create and version data assets in a workspace so that other members of your team can easily consume the data asset by using a name/version.

Example usage

Create data asset

To create a data asset, firstly define a data specification in a YAML file that provides a name, type and path for the data:

yml

```
# data-example.yml
$schema: https://azuremlschemas.azureedge.net/latest/data.schema.json
name: <name>
description: <description>
```

```
type: <type> # uri_file, uri_folder, mltable
path: https://<storage_name>.blob.core.windows.net/<container_name>/path
```

Then in the CLI, create the data asset:

Azure CLI

```
az ml data create --file data-example.yml --version 1
```

Dastore

An Azure Machine Learning datastore is a *reference* to an *existing* storage account on Azure. The benefits of creating and using a datastore are:

1. A common and easy-to-use API to interact with different storage types (Blob/Files /ADLS).
2. Easier to discover useful datastores when working as a team.
3. When using credential-based access (service principal/SAS/key), the connection information is secured so you don't have to code it in your scripts.

When you create a datastore with an existing storage account on Azure, you have the choice between two different authentication methods:

- **Credential-based** - authenticate access to the data using a service principal, shared access signature (SAS) token or account key. These credentials can be accessed by users who have *Reader* access to the workspace.
- **Identity-based** - authenticate access to the data using your Azure Active Directory identity or managed identity.

The table below summarizes which cloud-based storage services in Azure can be created as an Azure Machine Learning datastore and what authentication type can be used to access them.

Supported storage service	Credential-based authentication	Identity-based authentication
Azure Blob Container	✓	✓
Azure File Share	✓	

Supported storage service	Credential-based authentication	Identity-based authentication
Azure Data Lake Gen1	✓	✓
Azure Data Lake Gen2	✓	✓

ⓘ Note

The URI format to refer to a file/folder/mltable on a datastore is:

```
azureml://datastores/<name>/paths/<path>
```

MLTable

`mltable` is a way to abstract the schema definition for tabular data so that it is easier for consumers of the data to materialize the table into a Pandas/Dask/Spark dataframe.

💡 Tip

The ideal scenarios to use `mltable` are:

- The schema of your data is complex and/or changes frequently.
- You only need a subset of data (for example: a sample of rows or files, specific columns, etc).
- AutoML jobs requiring tabular data.

If your scenario does not fit the above then it is likely that URIs are a more suitable type.

A motivating example

Imagine a scenario where you have many text files in a folder:

text

```
|— my_data
|   |— file1.txt
|   |— file1_use_this.txt
```

```
|   └─ file2.txt  
|   └─ file2_use_this.txt  
.  
.  
.  
|   └─ file1000.txt  
|   └─ file1000_use_this.txt
```

Each text file has the following structure:

text

```
store_location date zip_code amount x y z noise_col1 noise_col2  
Seattle 20/04/2022 12324 123.4 true false true blah blah  
.  
.  
.  
London 20/04/2022 XX358YY 156 true true true blah blah
```

Some important features of this data are:

- The data of interest is only in files that have the following suffix: `_use_this.txt` and other file names that don't match should be ignored.
- The date should be represented as a date and not a string.
- The `x`, `y`, `z` columns are booleans, not strings.
- The store location is an index that is useful for generating subsets of data.
- The file is encoded in `ascii` format.
- Every file in the folder contains the same header.
- The first million records for `zip_code` are numeric but later on you can see they're alphanumeric.
- There are some dummy (noisy) columns in the data that aren't useful for machine learning.

You could materialize the above text files into a dataframe using Pandas and a URI:

Python

```
import glob  
import datetime  
import os  
import argparse  
import pandas as pd
```

```
parser = argparse.ArgumentParser()
parser.add_argument("--input_folder", type=str)
args = parser.parse_args()

path = os.path.join(args.input_folder, "*_use_this.txt")
files = glob.glob(path)

# create empty list
df1 = []

# dict of column types
col_types = {
    "zip": str,
    "date": datetime.date,
    "x": bool,
    "y": bool,
    "z": bool
}

# enumerate files into a list of dfs
for f in files:
    csv = pd.read_table(
        path=f,
        delimiter=" ",
        header=0,
        usecols=["store_location", "zip_code", "date", "amount", "x", "y",
"z"],
        dtype=col_types,
        encoding='ascii'
    )
    df1.append(csv)

# concatenate the list of dataframes
df = pd.concat(df1)
# set the index column
df.index_columns("store_location")
```

However, it will be the responsibility of the *consumer* of the data asset to parse the schema into a dataframe. In the scenario defined above, that means the consumers will need to independently ascertain the Python code to materialize the data into a dataframe.

Passing responsibility to the consumer of the data asset will cause problems when:

- **The schema changes (for example, a column name changes):** All consumers of the data must update their Python code independently. Other examples can be type changes, columns being added/removed, encoding change, etc.

- **The data size increases** - If the data gets too large for Pandas to process, then all the consumers of the data need to switch to a more scalable library (PySpark/Dask).

Under the above two conditions, `mltable` can help because it enables the creator of the data asset to define the schema in a single file and the consumers can materialize the data into a dataframe easily without needing to write Python code to parse the schema. For the above example, the creator of the data asset defines an `MLTable` file **in the same directory** as the data:

text

```
|— my_data
|   |— MLTable
|   |— file1.txt
|   |— file1_use_this.txt
|
.
.
.
```

The `MLTable` file has the following definition that specifies how the data should be processed into a dataframe:

YAML

```
type: mltable

paths:
  - pattern: ./*_use_this.txt

traits:
  - index_columns: store_location

transformations:
  - read_delimited:
      encoding: ascii
      header: all_files_same_headers
      delimiter: " "
  - keep_columns: ["store_location", "zip_code", "date", "amount", "x",
"y", "z"]
  - convert_column_types:
      - columns: ["x", "y", "z"]
        to_type: boolean
      - columns: "date"
        to_type: datetime
```

The consumers can read the data into dataframe using three lines of Python code:

Python

```
import mltable

tbl = mltable.load("./my_data")
df = tbl.to_pandas_dataframe()
```

If the schema of the data changes, then it can be updated in a single place (the MLTable file) rather than having to make code changes in multiple places.

Just like `uri_file` and `uri_folder`, you can create a data asset with `mltable` types.

Next steps

- [Install and set up the CLI \(v2\)](#)
- [Create datastores](#)
- [Create data assets](#)
- [Read and write data in a job](#)
- [Data administration](#)

